

BIMM 143 Lab 10

Ashley Martinez (PID: A17891957)

Table of contents

Candy Data	1
Favorite Candy	3
Overall Candy Rankings	7
Price Percent	10
Correlation Structure	12
Principal Component Analysis	13

As it is nearly Halloween and the half way point in the quarter, let's do a mini project to help us figure out the best candy!

Candy Data

Our data comes from the 538 website and is available as a CSV file:

```
candy<- read.csv("candy-data.txt",row.names=1)
flextable::flextable(head(candy,10))
```

chocolate	fruity	caramel	peanut	almond	nougat	crisped	rice	wafer	hard	bar	pluribus	s
1	0	1	0	0	0	1	0	0	0	1	0	0
1	0	0	0	0	1	0	0	0	0	1	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0
0	1	0	0	0	0	0	0	0	0	0	0	0
1	0	0	1	0	0	0	0	0	0	1	0	0

chocolate	fruity	caramel	peanut	almond	nougat	crisped	rice	wafer	hard	bar	pluribus	s
1	0	1	1	1	1	0	0	0	0	1	0	
0	0	0	1	0	0	0	0	0	0	0	1	
0	0	0	0	0	0	0	0	0	0	0	1	
0	1	1	0	0	0	0	0	0	0	0	0	

Q1. How many different candy types are in this dataset?

Answer: 85

```
nrow(candy)
```

[1] 85

```
candy |>
  nrow()
```

[1] 85

```
library(tidyverse)
```

```
-- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
v dplyr     1.1.4     v readr     2.1.5
vforcats    1.0.1     v stringr   1.5.2
v ggplot2   4.0.0     v tibble    3.3.0
v lubridate 1.9.4     v tidyverse 1.3.1
v purrr    1.1.0
-- Conflicts ----- tidyverse_conflicts() --
x dplyr::filter() masks stats::filter()
x dplyr::lag()    masks stats::lag()
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become non-conflicting
```

```
candy%>%
  nrow()
```

[1] 85

Q2. How many fruity candy types are in the dataset?

Answer: 38

```
sum(candy$fruity)
```

```
[1] 38
```

Favorite Candy

My favorite winpercent

```
candy["Twix",]$winpercent
```

```
[1] 81.64291
```

```
candy["Kit Kat",]$winpercent
```

```
[1] 76.7686
```

```
candy["Tootsie Roll Snack Bars",]$winpercent
```

```
[1] 49.6535
```

Q3. What is your favorite candy in the dataset and what is it's winpercent value?

Answer: 81.64%

Q4. What is the winpercent value for “Kit Kat”?

Answer: 76.77%

Q5. What is the winpercent value for “Tootsie Roll Snack Bars”?

Answer: 49.65%

Here is an alternative way to calculate winpercent.

```
library(dplyr)  
  
candy |>  
  filter(rownames(candy)=="Twix") |>  
  select(winpercent)
```

```

winpercent
Twix   81.64291

```

```

library("skimr")
skim(candy)

```

Table 2: Data summary

Name	candy
Number of rows	85
Number of columns	12
Column type frequency:	
numeric	12
Group variables	None

Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
chocolate	0	1	0.44	0.50	0.00	0.00	0.00	1.00	1.00	
fruity	0	1	0.45	0.50	0.00	0.00	0.00	1.00	1.00	
caramel	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
peanutyalmond	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
nougat	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
crispedricewafer	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
hard	0	1	0.18	0.38	0.00	0.00	0.00	0.00	1.00	
bar	0	1	0.25	0.43	0.00	0.00	0.00	0.00	1.00	
pluribus	0	1	0.52	0.50	0.00	0.00	1.00	1.00	1.00	
sugarpercent	0	1	0.48	0.28	0.01	0.22	0.47	0.73	0.99	
pricepercent	0	1	0.47	0.29	0.01	0.26	0.47	0.65	0.98	
winpercent	0	1	50.32	14.71	22.45	39.14	47.83	59.86	84.18	

Q6. Is there any variable/column that looks to be on a different scale to the majority of the other columns in the dataset?

Answer: The winpercent row seems to be on a different scale because all the other rows are on a scale from 0 to 1 while this row ranges from around 14 to 84, it is on a 0-100 scale.

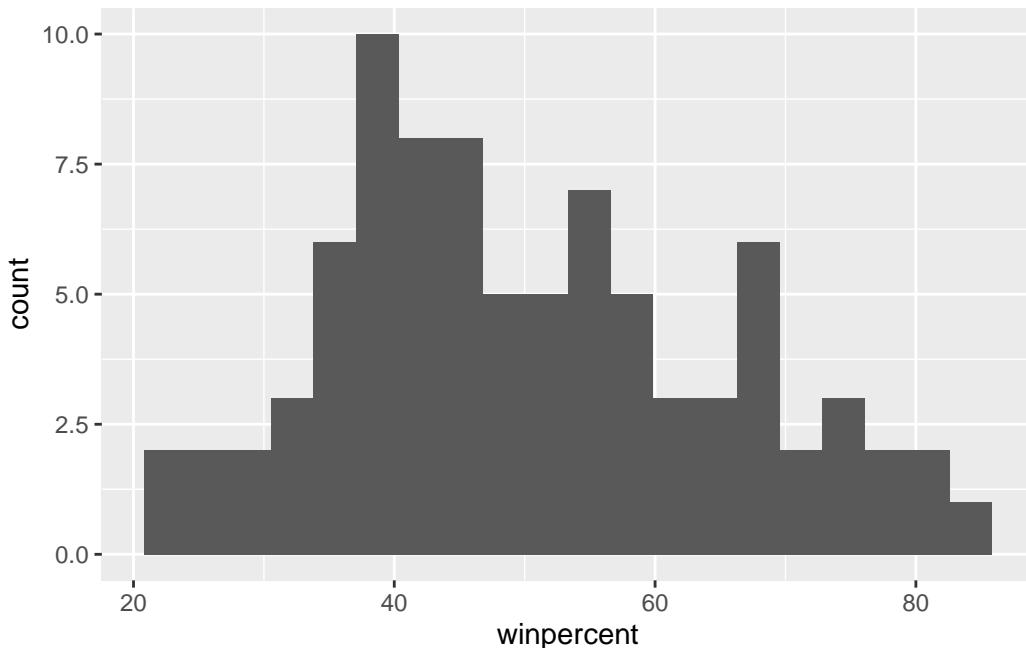
Q7. What do you think a zero and one represent for the candy\$chocolate column?

Answer: The candy doesn't contain chocolate if it is 0.

Q8. Plot a histogram of winpercent values

```
library(ggplot2)

ggplot(candy) +
  aes(winpercent) +
  geom_histogram(bins=20)
```



Q9. Is the distribution of winpercent values symmetrical?

Answer: The distribution is not symmetrical with the maximum skewed to the left.

Q10. Is the center of the distribution above or below 50%?

Answer: The center is below 50%.

```
mean(candy$winpercent)
```

[1] 50.31676

```
summary(candy$winpercent)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
22.45	39.14	47.83	50.32	59.86	84.18

Q11. On average is chocolate candy higher or lower ranked than fruit candy?

Answer: Chocolate is higher ranked.

To do this:

- #1. Find all chocolate candy in the dataset
- #2. Find their winpercent values
- #3. Calculate the mean of these values
- #4-6. Do the same for fruity candy
- #7. Compare mean winspercents of chocolate vs. fruity
- #8. Pick the highest as the winner

```
choc.inds<-candy$chocolate==1  
choc.win<-candy[choc.inds,]$winpercent  
choc.mean<-mean(choc.win)  
choc.mean
```

[1] 60.92153

```
fruit.inds<-candy$fruity==1  
fruit.win<-candy[fruit.inds,]$winpercent  
fruit.mean<-mean(fruit.win)  
fruit.mean
```

[1] 44.11974

```
mean(candy[candy$fruity==1,]$winpercent)
```

[1] 44.11974

Q12. Is this difference statistically significant?

Answer: Yes, the p value is less than 0.05.

```
t.test(choc.win,fruit.win)
```

Welch Two Sample t-test

```
data: choc.win and fruit.win
t = 6.2582, df = 68.882, p-value = 2.871e-08
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 11.44563 22.15795
sample estimates:
mean of x mean of y
 60.92153 44.11974
```

Overall Candy Rankings

Q13. What are the five least liked candy types in this set?

Answer: Nik L Nip, Boston Baked Beans, Chiclets, Super Bubble, and Jawbusters

```
candy |>
  arrange(winpercent) |>
  head(5)
```

	chocolate	fruity	caramel	peanut	yalmond	nougat		
Nik L Nip	0	1	0	0	0	0		
Boston Baked Beans	0	0	0	1	0	0		
Chiclets	0	1	0	0	0	0		
Super Bubble	0	1	0	0	0	0		
Jawbusters	0	1	0	0	0	0		
	crisped	rice	wafer	hard	bar	pluribus	sugarpercent	pricepercent
Nik L Nip	0	0	0	1	0	0.197	0.976	
Boston Baked Beans	0	0	0	1	0	0.313	0.511	
Chiclets	0	0	0	1	0	0.046	0.325	
Super Bubble	0	0	0	0	0	0.162	0.116	
Jawbusters	0	1	0	1	0	0.093	0.511	
	winpercent							
Nik L Nip	22.44534							
Boston Baked Beans	23.41782							
Chiclets	24.52499							
Super Bubble	27.30386							
Jawbusters	28.12744							

Q14. What are the top 5 all time favorite candy types out of this set?

Answer: Snickers, Kit Kat, Twix, Reese's Miniatures, and Reese's Peanut Butter cup

```
candy |>
  arrange(winpercent) |>
  tail(5)
```

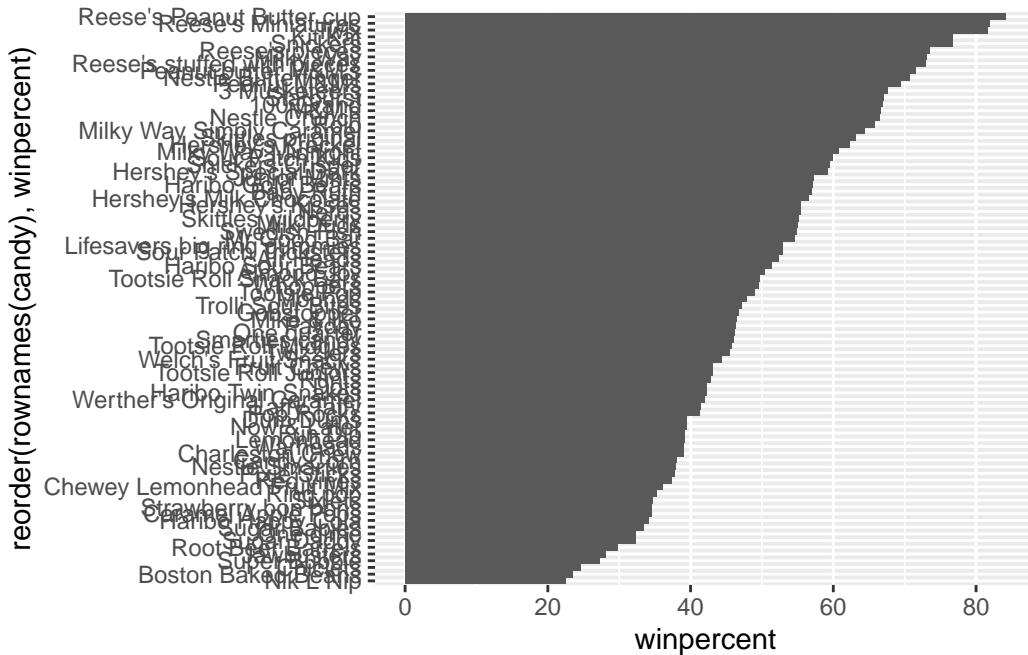
	chocolate	fruity	caramel	peanuty	almondy	nougat
Snickers	1	0	1		1	1
Kit Kat	1	0	0		0	0
Twix	1	0	1		0	0
Reese's Miniatures	1	0	0		1	0
Reese's Peanut Butter cup	1	0	0		1	0
	crispedrice	wafers	hard bar	pluribus	sugar	percent
Snickers	0	0	1	0	0	0.546
Kit Kat	1	0	1	0	0	0.313
Twix	1	0	1	0	0	0.546
Reese's Miniatures	0	0	0	0	0	0.034
Reese's Peanut Butter cup	0	0	0	0	0	0.720
	price	percent	winpercent			
Snickers	0.651	76.67378				
Kit Kat	0.511	76.76860				
Twix	0.906	81.64291				
Reese's Miniatures	0.279	81.86626				
Reese's Peanut Butter cup	0.651	84.18029				

Q15. Make a first barplot of candy ranking based on winpercent values.

Q16. This is quite ugly, use the reorder() function to get the bars sorted by winpercent?

```
library(ggplot2)

ggplot(candy) +
  aes(winpercent, reorder(rownames(candy),winpercent)) +
  geom_col()
```



Q17. What is the worst ranked chocolate candy?

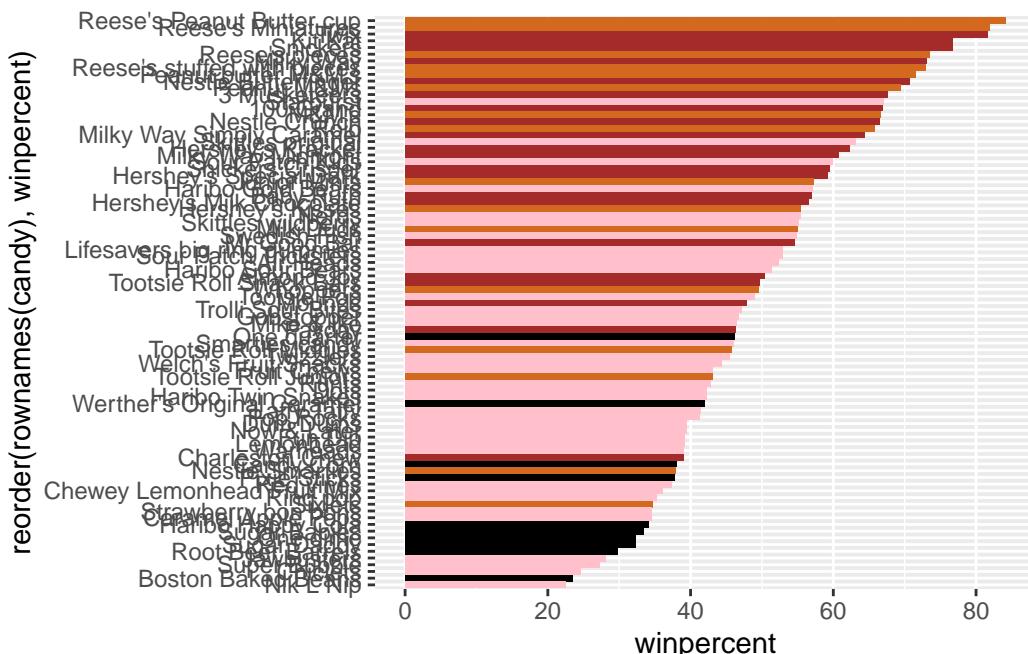
Answer: Sixlets

Q18. What is the best ranked fruity candy?

Answer: Starburst

```
my_cols<-rep("black", nrow(candy))
my_cols[as.logical(candy$chocolate)] = "chocolate"
my_cols[as.logical(candy$bar)] = "brown"
my_cols[as.logical(candy$fruity)] = "pink"

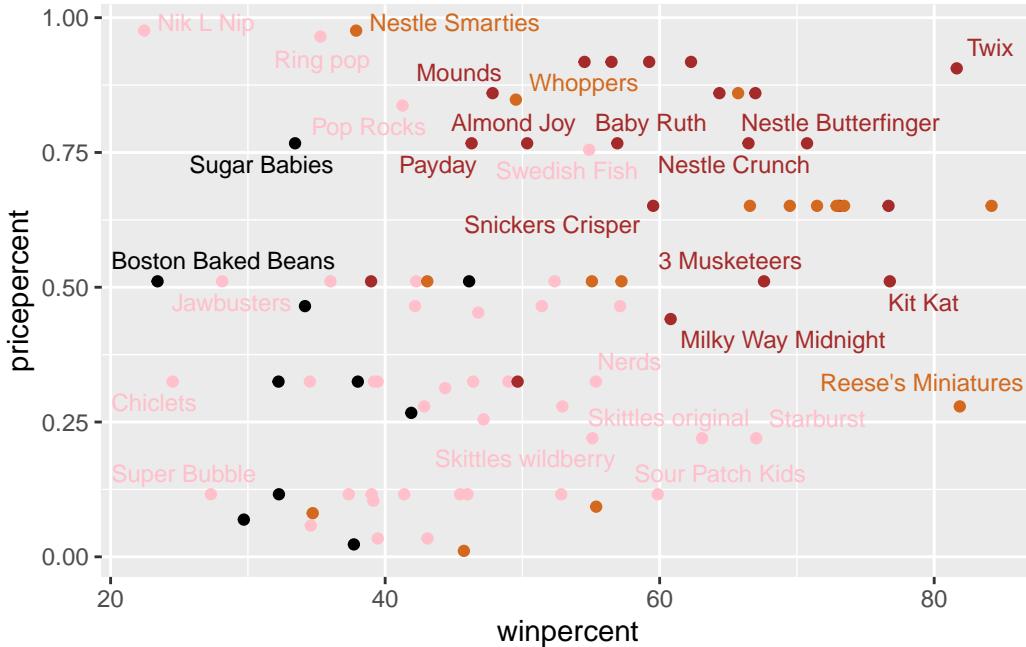
ggplot(candy) +
  aes(winpercent, reorder(rownames(candy),winpercent)) +
  geom_col(fill=my_cols)
```



Price Percent

```
library(ggrepel)
ggplot(candy) +
  aes(winpercent, pricepercent, label=rownames(candy)) +
  geom_point(col=my_cols) +
  geom_text_repel(col=my_cols, size=3.3, max.overlaps = 7)
```

Warning: ggrepel: 57 unlabeled data points (too many overlaps). Consider increasing max.overlaps



Q19. Which candy type is the highest ranked in terms of winpercent for the least money - i.e. offers the most bang for your buck?

Answer: Reese's Miniatures

Q20. What are the top 5 most expensive candy types in the dataset and of these which is the least popular?

Answer: Nik L Nip, Nestle Smarties, ring pop, Hershey's Krackel, Hershey's Milk Chocolate

```
ord <- order(candy$pricepercent, decreasing = TRUE)
head( candy[ord,c(11,12)], n=5 )
```

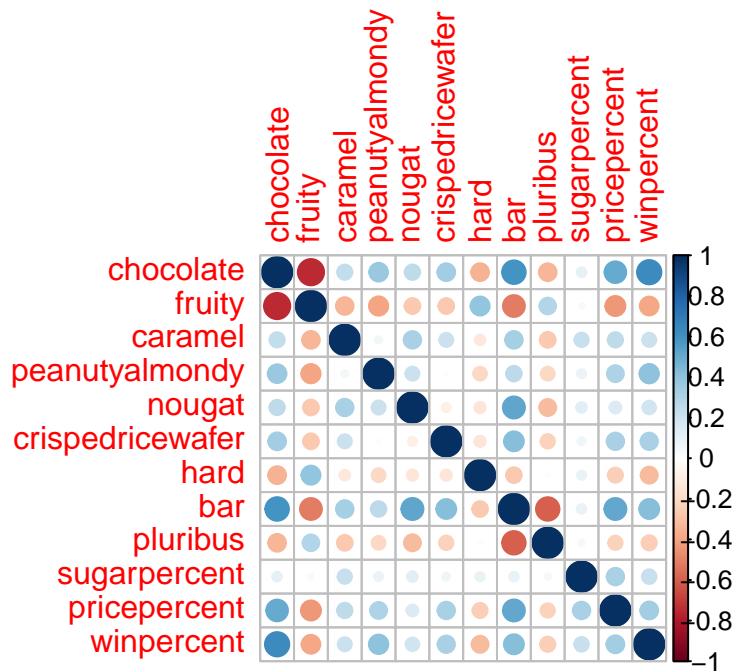
	pricepercent	winpercent
Nik L Nip	0.976	22.44534
Nestle Smarties	0.976	37.88719
Ring pop	0.965	35.29076
Hershey's Krackel	0.918	62.28448
Hershey's Milk Chocolate	0.918	56.49050

Correlation Structure

```
library(corrplot)
```

```
corrplot 0.95 loaded
```

```
cij<- cor(candy)
corrplot(cij)
```



Q22. Examining this plot what two variables are anti-correlated (i.e. have minus values)?

Answer: Chocolate and fruity

Q23. Similarly, what two variables are most positively correlated?

Answer: Chocolate and bar

Principal Component Analysis

The function to use is called ‘prcomp()’ with an optional ‘scale=T/F’ argument.

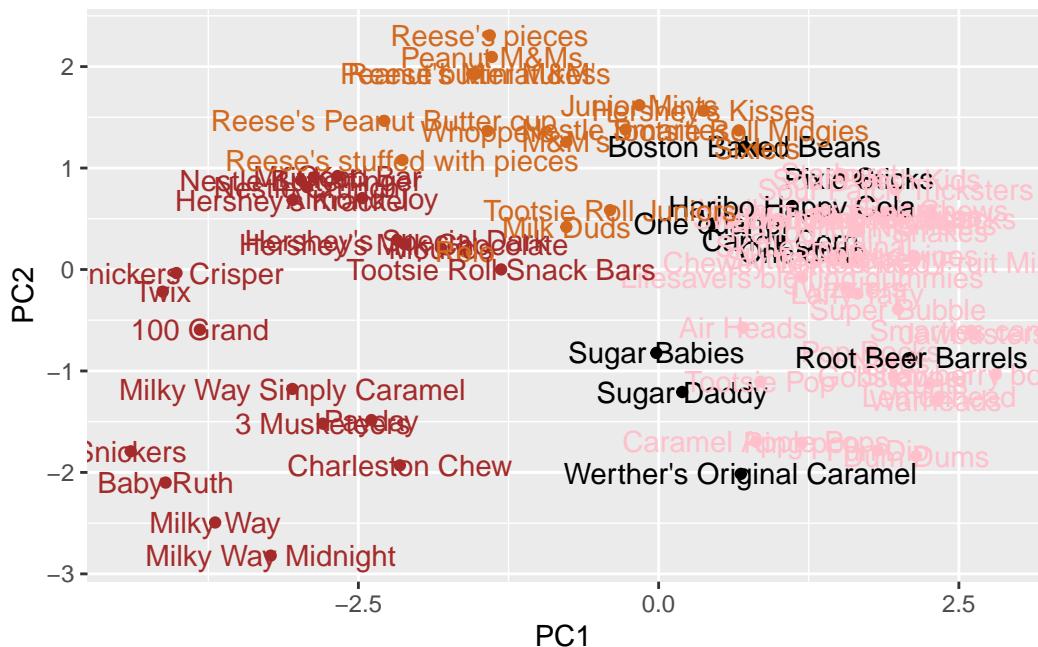
```
pca<- prcomp(candy, scale=TRUE)
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	2.0788	1.1378	1.1092	1.07533	0.9518	0.81923	0.81530
Proportion of Variance	0.3601	0.1079	0.1025	0.09636	0.0755	0.05593	0.05539
Cumulative Proportion	0.3601	0.4680	0.5705	0.66688	0.7424	0.79830	0.85369
	PC8	PC9	PC10	PC11	PC12		
Standard deviation	0.74530	0.67824	0.62349	0.43974	0.39760		
Proportion of Variance	0.04629	0.03833	0.03239	0.01611	0.01317		
Cumulative Proportion	0.89998	0.93832	0.97071	0.98683	1.00000		

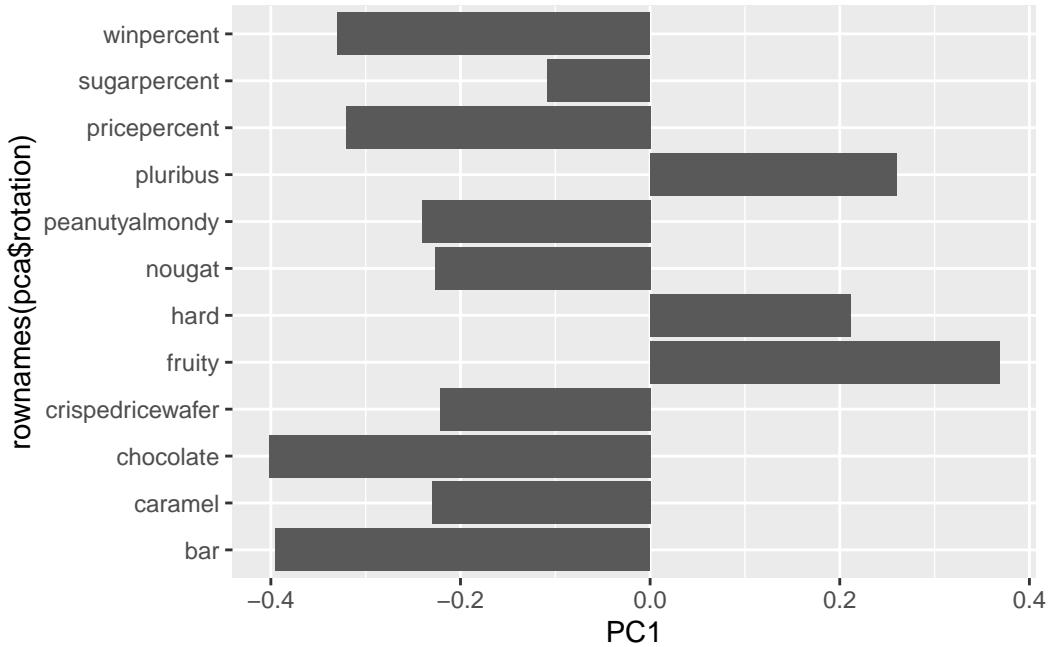
Our main PCA result figure

```
ggplot(pca$x)+
  aes(PC1,PC2,label=rownames(pca$x))+
  geom_point(col=my_cols)+
  geom_text(col=my_cols)
```



We should also examine the variable “loadings” or contributions of the original variable to the new PCs.

```
ggplot(pca$rotation)+  
  aes(PC1, rownames(pca$rotation))+  
  geom_col()
```

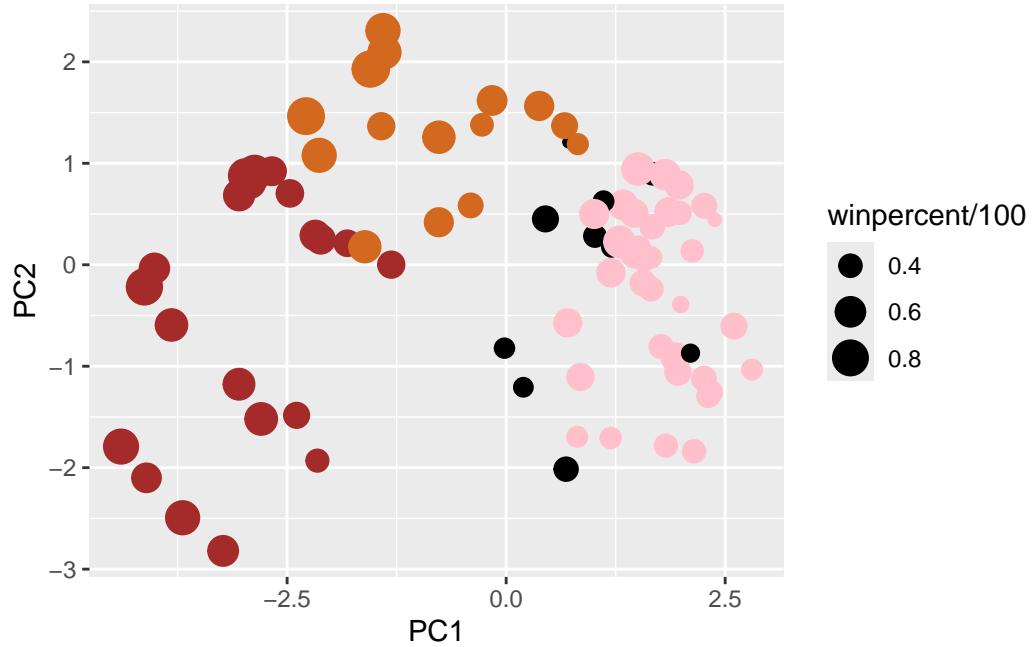


Q24. What original variables are picked up strongly by PC1 in the positive direction? Do these make sense to you?

Answer: Pluribus, hard, and fruity. These make sense because these tend to be the main trends for all fruity candy.

```
my_data <- cbind(candy, pca$x[,1:3])  
p <- ggplot(my_data) +  
  aes(x=PC1, y=PC2,  
      size=winpercent/100,  
      text=rownames(my_data),  
      label=rownames(my_data)) +  
  geom_point(col=my_cols)
```

p



Interactive plots that can be zoomed on and “brushed” over can be made with the **plotly** package. Its output is interactive and will not render to PDF.