

Class 12: RNA Seq Analysis

Ashley Martinez (PID: A17891957)

Table of contents

Background	1
Data Import	1
Toy with differential gene expression	3
Principal Component Analysis (PCA)	10
DESeq Analysis	12
Adding Annotation Data	14

Background

Today we will analyze some RNASeq data from Himes et. al. on the effects of a common steroid (dexamethasone) in airway smooth muscle cells (ASM cells).

A starting point is the “counts” data and “metadata” that contain the count values for each gene in their different experiments (i.e. cell lines with or without the drug).

Data Import

```
# Complete the missing code
counts <- read.csv("airway_scaledcounts.csv", row.names=1)
metadata <- read.csv("airway_metadata.csv")
```

Let's have a peek at these columns.

Q1. How many genes are in this dataset?

```
nrow(counts)
```

```
[1] 38694
```

```
nrow(metadata)
```

```
[1] 8
```

```
head(counts)
```

	SRR1039508	SRR1039509	SRR1039512	SRR1039513	SRR1039516
ENSG000000000003	723	486	904	445	1170
ENSG000000000005	0	0	0	0	0
ENSG000000000419	467	523	616	371	582
ENSG000000000457	347	258	364	237	318
ENSG000000000460	96	81	73	66	118
ENSG000000000938	0	0	1	0	2
	SRR1039517	SRR1039520	SRR1039521		
ENSG000000000003	1097	806	604		
ENSG000000000005	0	0	0		
ENSG000000000419	781	417	509		
ENSG000000000457	447	330	324		
ENSG000000000460	94	102	74		
ENSG000000000938	0	0	0		

```
head(metadata)
```

	id	dex	celltype	geo_id
1	SRR1039508	control	N61311	GSM1275862
2	SRR1039509	treated	N61311	GSM1275863
3	SRR1039512	control	N052611	GSM1275866
4	SRR1039513	treated	N052611	GSM1275867
5	SRR1039516	control	N080611	GSM1275870
6	SRR1039517	treated	N080611	GSM1275871

Q2. How many ‘control’ cell lines do we have?

Answer: 4

```
sum(metadata$dex == "control")
```

```
[1] 4
```

Toy with differential gene expression

To start, we will calculate the mean counts for all genes in the “control” experiments.

1. Extract all “control” columns from the ‘counts’ object
2. Calculate the mean for all rows (i.e. genes) of these “control” columns 3-4. Do the same for “treated”
3. Compare these ‘control.mean’ and ‘treated.mean’ values.

```
control inds<-metadata$dex=="control"  
control counts<-counts[,control inds]  
dim(control counts)
```

```
[1] 38694      4
```

```
control means<-rowMeans(control counts)
```

```
treated inds<-metadata$dex=="treated"  
treated counts<-counts[,treated inds]  
dim(treated counts)
```

```
[1] 38694      4
```

```
treated means<-rowMeans(treated counts)
```

Store the two control.means and treated.means together for bookkeeping.

```
meancounts<-data.frame(control means,treated means)  
head(meancounts)
```

	control.means	treated.means
ENSG000000000003	900.75	658.00
ENSG000000000005	0.00	0.00
ENSG000000000419	520.50	546.00
ENSG000000000457	339.75	316.50
ENSG000000000460	97.25	78.75
ENSG000000000938	0.75	0.00

Q3. How would you make the above code in either approach more robust? Is there a function that could help here?


Answer: rowMeans(counts[,metadata\$dex=="treated"])

Q4. Follow the same procedure for the treated samples (i.e. calculate the mean per gene across drug treated samples and assign to a labeled vector called treated.mean)

Answer: See treated.means above.

Q5a. Create a scatter plot of control vs treated.

```
plot(control.means,treated.means)
```



Q5 (b). You could also use the ggplot2 package to make this figure producing the plot below. What geom_?() function would you use for this plot?

Answer: Point


Q6. Try plotting both axes on a log scale. What is the argument to plot() that allows you to do this?

Make this plot a log plot to make the points more apparent.

```
plot(meancounts,log="xy")
```

```
Warning in xy.coords(x, y, xlabel, ylabel, log): 15032 x values <= 0 omitted  
from logarithmic plot
```

```
Warning in xy.coords(x, y, xlabel, ylabel, log): 15281 y values <= 0 omitted  
from logarithmic plot
```



We often talk about metrics like “log2 fold-change”

```
# treated/control  
log2(10/10)
```

```
[1] 0
```

The log2 can help show what changes are going up or down.

```
log2(10/40)
```

```
[1] -2
```

Let's calculate the log2 fold change for our treated over control mean counts.

```
meancounts$log2fc<-log2(meancounts$treated.means/
  meancounts$control.means)
head(meancounts)
```

	control.means	treated.means	log2fc
ENSG000000000003	900.75	658.00	-0.45303916
ENSG000000000005	0.00	0.00	NaN
ENSG00000000419	520.50	546.00	0.06900279
ENSG00000000457	339.75	316.50	-0.10226805
ENSG00000000460	97.25	78.75	-0.30441833
ENSG00000000938	0.75	0.00	-Inf

A common “rule of thumb” is a log2 fold change cutoff of +2 and -2 to call genes “Up regulated” or “Down regulated”.

```
zero.vals <- which(meancounts[,1:2]==0, arr.ind=TRUE)

to.rm <- unique(zero.vals[,1])
mycounts <- meancounts[-to.rm,]
head(mycounts)
```

	control.means	treated.means	log2fc
ENSG000000000003	900.75	658.00	-0.45303916
ENSG00000000419	520.50	546.00	0.06900279
ENSG00000000457	339.75	316.50	-0.10226805
ENSG00000000460	97.25	78.75	-0.30441833
ENSG00000000971	5219.00	6687.50	0.35769358
ENSG00000001036	2327.00	1785.75	-0.38194109

Q7. What is the purpose of the arr.ind argument in the which() function call above? Why would we then take the first column of the output and need to call the unique() function?

Answer: arr.ind is used equal to TRUE so results are given as a matrix and the which() function specifies columns 1:2 and identifies where the value is 0. Unique () makes sure no row is counted twice.

Now let's filter what genes are up or down regulated.

```
up.ind <- mycounts$log2fc > 2
down.ind <- mycounts$log2fc < (-2)

sum(up.ind)
```

```
[1] 250
```

```
sum(down.ind)
```

```
[1] 367
```

Q8. Using the up.ind vector above can you determine how many up regulated genes we have at the greater than 2 fc level?

Answer: 250

Q9. Using the down.ind vector above can you determine how many down regulated genes we have at the greater than 2 fc level?

Answer: 367

Q10. Do you trust these results? Why or why not?

Answer: Not necessarily because fold change can be large without being significant. DESeq2 will need to be used to determine significance.

```
library(DESeq2)
```

```
Loading required package: S4Vectors
```

```
Loading required package: stats4
```

```
Loading required package: BiocGenerics
```

```
Loading required package: generics
```

```
Attaching package: 'generics'
```

```
The following objects are masked from 'package:base':
```

```
as.difftime, as.factor, as.ordered, intersect, is.element, setdiff,  
setequal, union
```

```
Attaching package: 'BiocGenerics'
```

```
The following objects are masked from 'package:stats':
```

```
IQR, mad, sd, var, xtabs
```

```
The following objects are masked from 'package:base':
```

```
anyDuplicated, aperm, append, as.data.frame, basename, cbind,
colnames, dirname, do.call, duplicated, eval, evalq, Filter, Find,
get, grep, grepl, is.unsorted, lapply, Map, mapply, match, mget,
order, paste, pmax, pmax.int, pmin, pmin.int, Position, rank,
rbind, Reduce, rownames, sapply, saveRDS, table, tapply, unique,
unsplit, which.max, which.min
```

```
Attaching package: 'S4Vectors'
```

```
The following object is masked from 'package:utils':
```

```
findMatches
```

```
The following objects are masked from 'package:base':
```

```
expand.grid, I, unname
```

```
Loading required package: IRanges
```

```
Loading required package: GenomicRanges
```

```
Loading required package: Seqinfo
```

```
Loading required package: SummarizedExperiment
```

```
Loading required package: MatrixGenerics
```

```
Loading required package: matrixStats
```

```
Attaching package: 'MatrixGenerics'
```

```
The following objects are masked from 'package:matrixStats':
```

```
colAlls, colAnyNAs, colAnys, colAvgsPerRowSet, colCollapse,
colCounts, colCummaxs, colCummins, colCumprods, colCumsums,
colDiffs, colIQRDiffs, colIQRs, colLogSumExps, colMadDiffs,
colMads, colMaxs, colMeans2, colMedians, colMins, colOrderStats,
colProds, colQuantiles, colRanges, colRanks, colSdDiffs, colSds,
colSums2, colTabulates, colVarDiffs, colVars, colWeightedMads,
colWeightedMeans, colWeightedMedians, colWeightedSds,
colWeightedVars, rowAlls, rowAnyNAs, rowAnys, rowAvgsPerColSet,
rowCollapse, rowCounts, rowCummaxs, rowCummins, rowCumprods,
rowCumsums, rowDiffs, rowIQRDiffs, rowIQRs, rowLogSumExps,
rowMadDiffs, rowMads, rowMaxs, rowMeans2, rowMedians, rowMins,
rowOrderStats, rowProds, rowQuantiles, rowRanges, rowRanks,
rowSdDiffs, rowSds, rowSums2, rowTabulates, rowVarDiffs, rowVars,
rowWeightedMads, rowWeightedMeans, rowWeightedMedians,
rowWeightedSds, rowWeightedVars
```

```
Loading required package: Biobase
```

```
Welcome to Bioconductor
```

```
Vignettes contain introductory material; view with
'browseVignettes()'. To cite Bioconductor, see
'citation("Biobase")', and for packages 'citation("pkgname")'.
```

```
Attaching package: 'Biobase'
```

```
The following object is masked from 'package:MatrixGenerics':
```

```
rowMedians
```

```
The following objects are masked from 'package:matrixStats':
```

```
anyMissing, rowMedians
```

```
dds <- DESeqDataSetFromMatrix(countData=counts,
                               colData=metadata,
                               design=~dex)
```

```
converting counts to integer mode
```

```
Warning in DESeqDataSet(se, design = design, ignoreRank): some variables in  
design formula are characters, converting to factors
```

```
dds
```

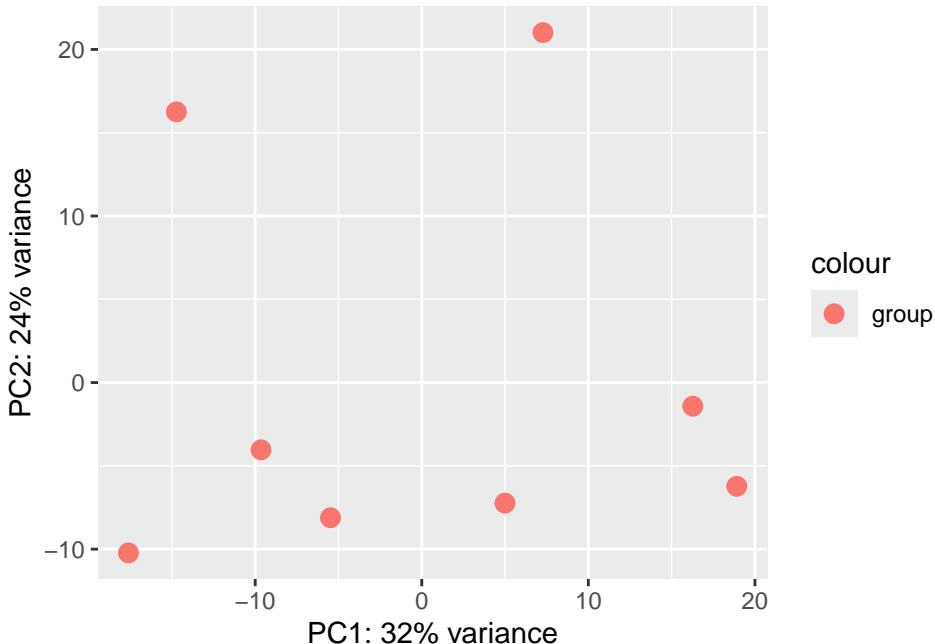
```
class: DESeqDataSet  
dim: 38694 8  
metadata(1): version  
assays(1): counts  
rownames(38694): ENSG00000000003 ENSG00000000005 ... ENSG00000283120  
    ENSG00000283123  
rowData names(0):  
colnames(8): SRR1039508 SRR1039509 ... SRR1039520 SRR1039521  
colData names(4): id dex celltype geo_id
```

Principal Component Analysis (PCA)

Now we will use PCA to analyze how the data samples are related to one another

```
vsd <- vst(dds, blind = FALSE)  
plotPCA(vsd, intgroup = c("dex"))
```

using ntop=500 top features by variance



```
pcaData <- plotPCA(vsd, intgroup=c("dex"), returnData=TRUE)
```

using ntop=500 top features by variance


```
head(pcaData)
```

	PC1	PC2	group	name	id	dex	celltype
SRR1039508	-17.607922	-10.225252	control	SRR1039508	SRR1039508	control	N61311
SRR1039509	4.996738	-7.238117	treated	SRR1039509	SRR1039509	treated	N61311
SRR1039512	-5.474456	-8.113993	control	SRR1039512	SRR1039512	control	N052611
SRR1039513	18.912974	-6.226041	treated	SRR1039513	SRR1039513	treated	N052611
SRR1039516	-14.729173	16.252000	control	SRR1039516	SRR1039516	control	N080611
SRR1039517	7.279863	21.008034	treated	SRR1039517	SRR1039517	treated	N080611
	geo_id	sizeFactor					
SRR1039508	GSM1275862	1.0193796					
SRR1039509	GSM1275863	0.9005653					
SRR1039512	GSM1275866	1.1784239					
SRR1039513	GSM1275867	0.6709854					
SRR1039516	GSM1275870	1.1731984					
SRR1039517	GSM1275871	1.3929361					

```

library(ggplot2)
percentVar <- round(100 * attr(pcaData, "percentVar"))
ggplot(pcaData) +
  aes(x = PC1, y = PC2, color = dex) +
  geom_point(size = 3) +
  xlab(paste0("PC1: ", percentVar[1], "% variance")) +
  ylab(paste0("PC2: ", percentVar[2], "% variance")) +
  coord_fixed() +
  theme_bw()

```



DESeq Analysis

Now we can run the DESeq analysis. We will need to reassign DESeq(dds) to dds.

```
dds <- DESeq(dds)
```

estimating size factors

estimating dispersions

gene-wise dispersion estimates

```
mean-dispersion relationship
```

```
final dispersion estimates
```

```
fitting model and testing
```

```
res <- results(dds)  
res
```

```
log2 fold change (MLE): dex treated vs control  
Wald test p-value: dex treated vs control  
DataFrame with 38694 rows and 6 columns  
  baseMean log2FoldChange    lfcSE      stat     pvalue  
  <numeric>      <numeric> <numeric> <numeric> <numeric>  
ENSG000000000003  747.1942   -0.350703  0.168242 -2.084514 0.0371134  
ENSG000000000005  0.0000     NA        NA        NA        NA  
ENSG000000000419  520.1342   0.206107  0.101042  2.039828 0.0413675  
ENSG000000000457  322.6648   0.024527  0.145134  0.168996 0.8658000  
ENSG000000000460  87.6826   -0.147143  0.256995 -0.572550 0.5669497  
...       ...       ...       ...       ...  
ENSG00000283115  0.000000     NA        NA        NA        NA  
ENSG00000283116  0.000000     NA        NA        NA        NA  
ENSG00000283119  0.000000     NA        NA        NA        NA  
ENSG00000283120  0.974916   -0.66825   1.69441  -0.394385 0.693297  
ENSG00000283123  0.000000     NA        NA        NA        NA  
  padj  
  <numeric>  
ENSG000000000003  0.163017  
ENSG000000000005  NA  
ENSG000000000419  0.175937  
ENSG000000000457  0.961682  
ENSG000000000460  0.815805  
...       ...  
ENSG00000283115  NA  
ENSG00000283116  NA  
ENSG00000283119  NA  
ENSG00000283120  NA  
ENSG00000283123  NA
```

Now lets make a summary of res with an adjusted cutoff.

```

res05 <- results(dds, alpha=0.05)
summary(res05)

out of 25258 with nonzero total read count
adjusted p-value < 0.05
LFC > 0 (up)      : 1237, 4.9%
LFC < 0 (down)     : 933, 3.7%
outliers [1]       : 142, 0.56%
low counts [2]      : 9033, 36%
(mean count < 6)
[1] see 'cooksCutoff' argument of ?results
[2] see 'independentFiltering' argument of ?results

```

Adding Annotation Data

We will now use AnnotationDbi to map between IDs.

```

library("AnnotationDbi")
library("org.Hs.eg.db")

```

```
columns(org.Hs.eg.db)
```

```

[1] "ACCNUM"        "ALIAS"          "ENSEMBL"         "ENSEMLPROT"    "ENSEMLTRANS"
[6] "ENTREZID"      "ENZYME"         "EVIDENCE"       "EVIDENCEALL"   "GENENAME"
[11] "GENETYPE"      "GO"              "GOALL"          "IPI"           "MAP"
[16] "OMIM"          "ONTOLOGY"       "ONTOLOGYALL"   "PATH"          "PFAM"
[21] "PMID"          "PROSITE"        "REFSEQ"         "SYMBOL"        "UCSCKG"
[26] "UNIPROT"

```

```

res$symbol <- mapIds(org.Hs.eg.db,
                      keys=row.names(res),
                      keytype="ENSEMBL",
                      column="SYMBOL",
                      multiVals="first")

```

```
'select()' returned 1:many mapping between keys and columns
```

```
head(res)
```

```
log2 fold change (MLE): dex treated vs control
Wald test p-value: dex treated vs control
DataFrame with 6 rows and 7 columns
  baseMean log2FoldChange      lfcSE      stat     pvalue
  <numeric>      <numeric> <numeric> <numeric> <numeric>
ENSG000000000003 747.194195      -0.350703  0.168242 -2.084514 0.0371134
ENSG000000000005   0.000000          NA        NA        NA        NA
ENSG00000000419  520.134160      0.206107  0.101042  2.039828 0.0413675
ENSG00000000457  322.664844      0.024527  0.145134  0.168996 0.8658000
ENSG00000000460   87.682625      -0.147143  0.256995 -0.572550 0.5669497
ENSG00000000938   0.319167      -1.732289  3.493601 -0.495846 0.6200029
  padj      symbol
  <numeric> <character>
ENSG000000000003  0.163017      TSPAN6
ENSG000000000005    NA        TNMD
ENSG00000000419   0.175937      DPM1
ENSG00000000457   0.961682      SCYL3
ENSG00000000460   0.815805      FIRRM
ENSG00000000938    NA        FGR
```

Q11. Run the mapIds() function two more times to add the Entrez ID and UniProt accession and GENENAME as new columns called res\$entrez, res\$uniprot and res\$genename.

Answer:

```
res$entrez <- mapIds(org.Hs.eg.db,
                      keys=row.names(res),
                      column="ENTREZID",
                      keytype="ENSEMBL",
                      multiVals="first")
```

```
'select()' returned 1:many mapping between keys and columns
```

```
res$uniprot <- mapIds(org.Hs.eg.db,
                      keys=row.names(res),
                      column="UNIPROT",
                      keytype="ENSEMBL",
                      multiVals="first")
```

```
'select()' returned 1:many mapping between keys and columns
```

```
res$genename <- mapIds(org.Hs.eg.db,
                      keys=row.names(res),
                      column="GENENAME",
                      keytype="ENSEMBL",
                      multiVals="first")
```

```
'select()' returned 1:many mapping between keys and columns
```

```
head(res)
```

```
log2 fold change (MLE): dex treated vs control
Wald test p-value: dex treated vs control
DataFrame with 6 rows and 10 columns
      baseMean log2FoldChange     lfcSE      stat    pvalue
      <numeric>     <numeric> <numeric> <numeric> <numeric>
ENSG000000000003 747.194195     -0.350703  0.168242 -2.084514 0.0371134
ENSG000000000005  0.000000        NA        NA        NA        NA
ENSG000000000419 520.134160      0.206107  0.101042  2.039828 0.0413675
ENSG000000000457 322.664844      0.024527  0.145134  0.168996 0.8658000
ENSG000000000460  87.682625     -0.147143  0.256995 -0.572550 0.5669497
ENSG000000000938  0.319167     -1.732289  3.493601 -0.495846 0.6200029
      padj      symbol     entrez      uniprot
      <numeric> <character> <character> <character>
ENSG000000000003 0.163017      TSPAN6      7105 AOA087WYV6
ENSG000000000005   NA          TNMD      64102 Q9H2S6
ENSG000000000419 0.175937      DPM1       8813 H0Y368
ENSG000000000457 0.961682      SCYL3      57147 X6RHX1
ENSG000000000460 0.815805      FIRRM      55732 A6NFP1
ENSG000000000938   NA          FGR       2268 B7Z6W7
      genename
      <character>
ENSG000000000003      tetraspanin 6
ENSG000000000005      tenomodulin
ENSG000000000419 dolichyl-phosphate m..
ENSG000000000457 SCY1 like pseudokina..
ENSG000000000460 FIGNL1 interacting r..
ENSG000000000938 FGR proto-oncogene, ..
```